

Structured Metacognition via 2-Level GRPO-Embedded Monte Carlo Tree Search for Decoding Policy Optimization

Niel Ok
Stanford University
nielok@stanford.edu

May 2025

Abstract

We introduce a general and scalable test-time inference framework for text diffusion models that performs structured search over decoding policies independently for each input prompt rather than decoding output sequences with a fixed strategy. Our approach embeds Group-Relative Policy Optimization (GRPO) within Monte Carlo Tree Search (MCTS), where each node represents a partial decoding strategy parameterized by fixed-length schedules over temperature, remasking strategy, and compute allocation. For each prompt, GRPO-MCTS constructs a search tree: before expanding a node, it conducts exploratory rollouts using multiple sampled children and applies GRPO to refine the node’s sampling logits. These updates occur across sibling children (same tree depth) and are tailored to the specific prompt, enabling stable advantage estimation under reward sparsity. Final children are sampled from the optimized logits, allowing the tree to grow in directions aligned with high-reward decoding behaviors. This forms a two-level structure: MCTS searches vertically over inference schedules, while GRPO optimizes locally across policies. Scaling is controlled by the breadth of tree search and the precision of GRPO updates. We frame this as structured metacognition, learning how to think, and demonstrate its effectiveness on the FO-LIO logical entailment benchmark.

We argue that decoding should be adaptive, introspective, and structurally optimized, even at test time, and that text diffusion models give us the fine-grained control we need to do so.

We introduce **GRPO-MCTS**, a decoding-time algorithm for text diffusion models that performs structured search over decoding policies independently for each input prompt. GRPO-MCTS uses Monte Carlo Tree Search (MCTS) [1] to explore the space of decoding strategies, where each node encodes a partial policy defined by temperature schedules, remasking heuristics, block lengths, and compute allocations. A full rollout defines a complete decoding strategy that guides text generation for that specific prompt.

To optimize these strategies, we embed a two-level variant of Group-Relative Policy Optimization (GRPO) [4] at every node. Before sampling children, each node does random rollouts and applies GRPO updates using evaluations on the current prompt. This allows refining of internal logits, pre-emptively informing structural search decisions. GRPO is conducted (1) across sibling children at the same tree depth, and (2) within the scope of a single prompt, stabilizing advantage estimation and enabling policy learning even under sparse or noisy rewards. As a result, GRPO-MCTS adaptively co-evolves its search and policy in real time—learning not just what to say, but how to think, one prompt at a time.

1 Introduction

Language model decoding is traditionally governed by static, globally fixed strategies such as greedy decoding, temperature sampling, or nucleus sampling. These methods overlook a central fact: generation quality depends not only on what a model decodes, but how it decodes.

2 Decoding Policies as Structured Programs

Let x be a prompt and $y = (y_1, \dots, y_T)$ a generated sequence. We define a decoding policy π_s that governs the model’s inference behavior through a structured schedule $s = \{(t_i, r_i, b_i, e_i)\}_{i=1}^N$, where N is the (fixed) number

of decoding blocks held constant across all policies for a given experiment and:

- t_i is the sampling temperature applied in block i ,
- r_i is the remasking strategy used in block i (e.g., low-confidence, random),
- b_i is the number of base generation steps assigned to block i ,
- e_i is the proportion of residual steps (adaptive compute) allocated to block i .

This schedule specifies not only local sampling behavior, but also a global strategy for allocating compute across the generation process. Each decoding policy can thus be interpreted as a structured program: it dynamically controls how the model allocates attention, uncertainty, and effort during generation.

Sampling from π_s generates an output $y \sim \pi_s(\cdot | x)$, which is evaluated by a reward function $R(y)$. Our objective is to identify high-performing schedules for each prompt:

$$s^* = \arg \max_s \mathbb{E}_{y \sim \pi_s} [R(y)]. \quad (1)$$

In our framework, this search is implemented via Monte Carlo Tree Search (MCTS), where each node represents a prefix of s , and full rollouts define complete decoding policies. These are iteratively refined using Group-Relative Policy Optimization (GRPO), enabling adaptive inference at test time.

3 MCTS over Partial Decoding Policies

We construct a tree \mathcal{T} in which each node n_d at depth d represents a partial decoding policy $s_{\leq d}$ —a prefix of a structured schedule. Each schedule element at depth d is a tuple (t_d, r_d, b_d, e_d) corresponding to temperature, remasking strategy, base steps, and extra step proportion. A complete decoding policy s is produced by rolling out a path from the root to a leaf of the tree. The tree has a maximum depth of N , corresponding to the fixed number of blocks in each decoding policy. Each full rollout thus represents a complete blockwise schedule of length N .

MCTS in this context explores structural allocation policies: how to segment the generation process into blocks, how to distribute base and residual compute across those blocks, and which local sampling heuristics to apply. These decisions represent architectural control over the inference process, not just token-level behavior.

Before expanding a node, GRPO-MCTS performs a *Phase 1 update*: R exploratory child policies are sampled from the node’s current temperature and remasking logits. These are randomly rolled out into full decoding policies, and their outputs are evaluated via a prompt-specific reward function $R(y)$. The reward signals are used in a two-level Group-Relative Policy Optimization (GRPO) procedure:

- *Level 1 (horizontal)*: GRPO is applied across sibling children at the same depth, optimizing logits based on relative performance.
- *Level 2 (prompt-wise)*: Since the algorithm is applied per prompt, GRPO is applied independently for each prompt, allowing for adjustment to prompt-specific difficulty.

After GRPO updates the node’s logits, k final children are sampled from the improved distribution and added to the tree, and rewards are normalized and propagated up the tree based on those final children. This coupling of search and local learning ensures that structural expansions reflect reward-aligned behaviors.

Over successive iterations, MCTS grows the tree along promising directions, while GRPO sharpens the sampling behavior within each node. This co-evolution enables the model to discover and execute high-performing inference procedures tailored to individual prompts.

4 Two-Level GRPO for Fair and Structured Credit Assignment

To refine decoding behavior during tree expansion, we apply Group-Relative Policy Optimization (GRPO) across a group of random, exploratory rollouts. Let x be the current prompt and $\mathcal{G}_d = \{s_1, \dots, s_k\}$ the group of sampled child policies at a node of depth d .

Each policy $s_j \in \mathcal{G}_d$ is rolled out to produce an output $y_j \sim \pi_{s_j}(\cdot | x)$, which is evaluated by a scalar reward $R_j = R(y_j)$. We then apply GRPO to compute a gradient objective for prompt x :

$$\mathcal{L}_{\text{GRPO}} = - \sum_{j=1}^k (R_j - \bar{R}_{\mathcal{G}_d}) \log P(s_j), \quad (2)$$

where $\bar{R}_{\mathcal{G}_d} = \frac{1}{k} \sum_{j=1}^k R_j$ is the average reward across the group for prompt x , and $P(s_j)$ is the sampling probability of policy s_j under the node’s current logits.

After computing advantages, we aggregate the advantages for each policy and normalize them via z-score over

the group \mathcal{G}_d , producing group-normalized advantages:

$$A_j = \frac{1}{N} \sum_{i=1}^N (R_j - \bar{R}_{\mathcal{G}_d}) \quad (3)$$

$$\bar{A} = \frac{1}{k} \sum_{j=1}^k A_j \quad (4)$$

$$\sigma_A^2 = \frac{1}{k} \sum_{j=1}^k (A_j - \bar{A})^2 \quad (5)$$

$$\hat{A}_j = \frac{A_j - \bar{A}}{\sigma_A + \epsilon} \quad (6)$$

These advantages \hat{A}_j are then used to compute policy gradients and update the temperature and remasking logits via backpropagation.

This two-level formulation ensures:

- *Prompt-specific optimization*: The optimization objective is tailored to the given prompt.
- *Structure-level relativity*: Credit assignment is local to each node’s group of children, ensuring refinement focuses on intra-structural distinctions.

5 Algorithm Overview

We summarize the GRPO-MCTS inference-time procedure in Algorithm 1. The core idea is to search not just over output sequences but over decoding strategies themselves. Each node in the MCTS tree represents a partial decoding policy, specifically, a prefix of block-level parameters including temperature, remasking strategy, step allocation, and extra compute proportion. As the tree grows, full decoding policies are constructed by rolling out paths from the root to leaves.

To expand a node, GRPO-MCTS samples multiple exploratory children, rolls them out, and evaluates their outputs using a prompt-specific reward function. Group-Relative Policy Optimization (GRPO) is then applied to refine the node’s sampling logits based on prompt-normalized, z-score-adjusted advantages. After this update, a fixed number of high-reward children are sampled and added to the tree. These steps are repeated recursively, enabling the tree to grow in directions aligned with effective inference behavior. GRPO-MCTS operates independently for each prompt and requires no modification of the underlying language model, making it a general-purpose method for decoding-time optimization.

Algorithm 1 GRPO-MCTS for Decoding Policy Optimization (per prompt)

```

1: procedure GRPO_MCTS( $M, T, x, \ell, S, I, R, b, k$ )
2:   Initialize root node  $r$  with empty
   DecodingPolicyState
3:   Initialize optimizer  $\mathcal{O}$  over logits in subtree of  $r$ 
4:    $\mathcal{K} \leftarrow$  IDs of all logits tracked by  $\mathcal{O}$ 
5:   for  $i = 1$  to  $I$  do
6:      $n \leftarrow r$ 
7:     while  $n$  is fully expanded and not terminal do
8:        $n \leftarrow \text{BestChild}(n)$ 
9:     end while
10:    if  $n$  is not terminal then
11:      Phase 1: GRPO Pretraining Step
12:       $\mathcal{G} \leftarrow R$  exploratory children sampled
      from logits at  $n$ 
13:      for all  $s_j \in \mathcal{G}$  do
14:         $\pi_j \leftarrow \text{RolloutPolicy}(s_j, S)$ 
15:         $y_j \leftarrow M(x, \pi_j)$ 
16:         $r_j \leftarrow R(y_j, \ell)$ 
17:      end for
18:      Compute z-scored advantages  $\hat{A}_j$  across
      group  $\mathcal{G}$ 
19:      Apply GRPO update:  $\nabla_{\theta} \mathcal{L}_{\text{GRPO}} =$ 
       $-\sum_j \hat{A}_j \nabla_{\theta} \log P(s_j)$ 
20:      Update logits at node  $n$  using  $\mathcal{O}$ 
21:      Phase 2: Final Sampling and Value
      Propagation
22:       $\mathcal{C} \leftarrow b$  children sampled from updated
      logits at  $n$ 
23:      for all  $c \in \mathcal{C}$  do
24:         $\pi_c \leftarrow \text{RolloutPolicy}(c, S)$ 
25:         $y_c \leftarrow M(x, \pi_c)$ 
26:         $r_c \leftarrow R(y_c, \ell)$ 
27:      end for
28:      Compute z-scored advantages  $\hat{A}_c$  across  $\mathcal{C}$ 
29:      for all  $c \in \mathcal{C}$  do
30:        Propagate  $\hat{A}_c$  upward from  $c$  to root  $r$ 
31:      end for
32:      Add any new logits in  $\mathcal{C}$  to  $\mathcal{O}$  if not in  $\mathcal{K}$ 
33:    end if
34:  end for
35:   $\mathcal{L}_{\text{leaf}} \leftarrow$  all leaves in tree rooted at  $r$ 
36:  Sort  $\mathcal{L}_{\text{leaf}}$  by average value (value_sum / visits)
37:  return Decoding policies from top- $k$  leaves
38: end procedure

```

6 Theoretical Analysis

We analyze convergence properties of two-level Group-Relative Policy Optimization (GRPO) in the context of structured decoding policy optimization. Let \mathcal{P}_d denote the set of full decoding policies that share a common prefix of length d —i.e., policies generated from a single parent node at depth d in the MCTS tree. Let x denote the prompt, and let $R(y)$ be a bounded scalar reward function.

Our goal is to maximize the expected return for the prompt, defined as:

$$J(s) = \mathbb{E}_{y \sim \pi_s(\cdot|x)}[R(y)] - \bar{R}_{\mathcal{P}_d}, \quad (7)$$

where $\bar{R}_{\mathcal{P}_d} = \frac{1}{|\mathcal{P}_d|} \sum_{s' \in \mathcal{P}_d} \mathbb{E}_{y \sim \pi_{s'}(\cdot|x)}[R(y)]$ is the mean reward over the sibling group \mathcal{P}_d for prompt x .

6.1 Assumptions

We adopt the following standard assumptions to support convergence:

1. **Bounded Rewards:** $R(y) \in [0, R_{\max}]$ for all outputs y .
2. **Smooth Policies:** The sampling distribution π_s is differentiable and Lipschitz-continuous in the scheduling parameters of s (e.g., logits over temperature/remasking options).
3. **Finite Action Space:** The set of possible block parameters (temperature, remasking, block length, and extra step proportion) is finite.
4. **Persistent Exploration:** The UCB-based node selection ensures every policy $s \in \mathcal{P}_d$ is visited infinitely often.
5. **Log-Probability Gradients:** The sampling log-probability $\log P(s)$ is differentiable and gradients $\nabla \log P(s)$ are accessible via softmax logits.
6. **Fixed Block Count:** Each decoding policy consists of exactly N blocks, where N is constant for a given experiment.

6.2 Gradient-Based Convergence

Let $P(s)$ denote the sampling probability of decoding policy $s \in \mathcal{P}_d$ from the node’s softmax distribution over child options. For the prompt x , the GRPO objective is defined as:

$$\mathcal{L}_{\text{GRPO}} = - \sum_{s \in \mathcal{P}_d} \hat{A}(s) \log P(s), \quad (8)$$

where $R(s) = R(\pi_s(x))$ is the observed reward for prompt x under policy s .

Taking the gradient with respect to the policy parameters (i.e., logits underlying $P(s)$), we obtain the standard log-likelihood trick:

$$\nabla \mathcal{L}_{\text{GRPO}} = - \sum_{s \in \mathcal{P}_d} \hat{A}(s) \nabla \log P(s). \quad (9)$$

Sampling-based estimation yields an unbiased stochastic policy gradient:

$$\nabla \mathcal{L}_{\text{GRPO}} = \mathbb{E}_{s \sim P(\cdot)} [(R(s) - \bar{R}_{\mathcal{P}_d}) \nabla \log P(s)]. \quad (10)$$

This is all for one prompt, hence two-level GRPO, optimizing for same-depth nodes and the same prompt. This gradient is an unbiased estimator of $\nabla J(s)$, and under Robbins-Monro conditions for stochastic approximation and the assumptions above, the GRPO update converges almost surely to a local optimum of the relative expected return $J(s)$ within each structural group \mathcal{P}_d .

6.3 Interpretation

The two-level GRPO design decouples prompt-level normalization from structure-level competition, enabling stable optimization in sparse or noisy reward settings. By normalizing rewards across policies per prompt, the algorithm ensures that the policy is optimized specifically for the given prompt. Simultaneously, by comparing policies only within structurally consistent sibling groups \mathcal{P}_d , GRPO ensures that the refinement signal is locally aligned to meaningful architectural choices.

This structure-aware credit assignment enables the model to learn prompt-specific decoding behaviors without conflating unrelated inference strategies. Provided that MCTS explores sufficiently, ensuring each policy is visited infinitely often, two-level GRPO converges to local optima that are both prompt-fair and structure-aware, making it well-suited for test-time reasoning tasks where generalization must be adapted per-instance.

7 Experimental Setup

We evaluate GRPO-MCTS on logical reasoning tasks using the FOLIO dataset [2], which consists of multi-premise entailment questions with sparse supervision and nuanced semantic structure. We use LLaDA [3] as our text diffusion model, but any text diffusion model can be used. Each decoding policy consists of a blockwise schedule over the following discrete action space (with a fixed number of blocks N):

- **Temperature choices:**
 $t_i \in \{0.0, 0.1, 0.2\}$,
- **Remasking strategy choices:**
 $r_i \in \{\text{low-confidence}, \text{random}\}$,
- **Block lengths:**
 $b_i \in \{1, \dots, B\}$, subject to $\sum_{i=1}^N b_i = T$,
- **Extra step proportions:**
 $e_i \in [0, 1]$, with $\sum_{i=1}^N e_i \leq 1$.

Each candidate policy is evaluated on a shared prompt and reward is computed using an external verifier model. GRPO is applied at two levels: (1) across sibling policies at the same tree depth, and (2) normalized per prompt to ensure fairness and stability. All experiments use the 8B instruct-tuned variant of LLaDA with float16 inference and are run per prompt to reflect test-time usage.

Scaling in our framework is governed by three key factors. First, the *branching factor* determines how many children are sampled at each node, shaping the breadth of exploration as the decoding policy state is incrementally constructed. Second, the *number of Phase 1 groups* controls how many independent sets of rollouts are used to perform the GRPO update. Each group provides a localized comparison across sibling policies, enabling more accurate updates to sampling logits such as temperature and remasking strategy. Third, the *number of rollouts per group* affects how well each group’s update is tuned. A higher count leads to more stable gradient estimates and more refined policy adjustments. Together, these parameters control the overall trade-off between search breadth, update quality, and inference cost.

8 Research Prototype Results

We evaluate the effectiveness of GRPO-embedded MCTS on test-time logical reasoning using prompts from the FO-LIO dataset. Each prompt contains a set of facts followed by a question requiring compositional inference. Decoding policies sampled by our framework guide the generation process, and outputs are evaluated using an LLM-as-a-Judge for correctness, validity, and coherence. We set the number of blocks to 4.

8.1 Prompt and Trajectory Example

The following is an example of a prompt used in our evaluation:

The following facts are given: Phoneix’s music is classified under the indie pop genre. is a band from

France. French bands write songs in French or in English. Aside from indie pop, pop rock and synth-pop are two other genres of music. Phoenix has no songs in French.

Question: Is the conclusion ”Phoneix’s music is classified under the pop rock genre.” logically entailed by the above facts? Think out loud step by step and, on a new line, write one of the following three options by itself as your final answer: **False**, **Uncertain**, or **True**.

During inference, our framework gets the model to explore a structured space of decoding policies. Each policy defines a schedule over temperature, remasking strategy, block length, and extra step proportions. Below is one such decoding policy that was discovered by test-time reinforcement learning using our GRPO-embedded MCTS framework for the prompt above:

Decoding Policy (`test_time_policy_0`):

- Temperature schedule:
[0.2, 0.2, 0.2, 0.2]
- Remasking strategy schedule:
[random, low_confidence,
low_confidence, low_confidence]
- Block schedule:
[37, 58, 26, 7]
- Extra step proportions:
[0.24, 0.38, 0.04, 0.34]

This policy yielded a successful reasoning trajectory that logically entailed the correct answer:

Model Output:

Let’s break down the facts: 1. Phoneix’s music is classified under the indie pop genre. 2. Phoenix is a band from France. 3. French bands write songs in French or in English. 4. Aside from indie pop, pop rock and synth-pop are two other genres of music. 5. Phoenix has no songs in French. Given these facts, the conclusion ”Phoneix’s music is classified under the pop rock genre” is not logically entailed. Phoneix’s music is definit as classified under the indie pop genre. So, the final answer is:

False

This example demonstrates the ability of GRPO-MCTS to discover effective decoding strategies that enable consistent reasoning under sparse supervision.

8.2 Links to Results

We provide the following artifacts from our research prototype to enable replication and inspection of generated policies and reasoning trajectories:

- **Metadata:** High-level metadata describing parameters for the GRPO-embedded MCTS algorithm, each sampled decoding policy, and its associated reward.
 - Metadata JSON
- **Tree Snapshot:** Full serialized tree snapshot containing decoding policies explored by GRPO-MCTS, including all sampled children, rewards, logits, and value estimates.
 - Tree Snapshot
- **Verbose Log:** Full prompt-level rollout trace, including reasoning outputs and reward assignments for each policy.
 - Output Log

9 Applications and Outlook

- **Prompt-specific decoding optimization:** GRPO-MCTS dynamically adapts inference behavior to the specific reasoning demands of each prompt, rather than relying on globally fixed decoding strategies.
- **Verifier-guided generation:** The framework is agnostic to the reward source and can be paired with external verifiers or task-specific reward models.
- **Policy distillation:** Top-performing decoding policies identified by GRPO-MCTS can be distilled into lightweight student decoders or used to guide future model pretraining.
- **Exploration extensions:** Future work may incorporate predictive entropy or calibration metrics into MCTS exploration bonuses, or extend the GRPO policy space to continuous actions. Multi-prompt optimization for similar kinds of prompts may also be explored.

10 Conclusion

We introduce 2-level GRPO-MCTS, a test-time decoding optimization framework that jointly performs structured search and credit assignment over blockwise decoding policies. By embedding Group-Relative Policy Optimization within Monte Carlo Tree Search, the method

adapts inference behavior to each input prompt via competitive local learning and structural exploration. This enables adaptive generation, prompt-fair credit assignment, and a novel form of structured metacognition in autoregressive language modeling. GRPO-MCTS is modular, gradient-free with respect to the base model, and compatible with any sampling-based decoder, offering a general-purpose foundation for test-time inference optimization.

11 Acknowledgements

We gratefully acknowledge compute support provided by researchers at DeepMind, in preparation for DeepMind’s Gemini Diffusion release.

References

- [1] Rémi Coulom. Efficient selectivity and backup operators in monte-carlo tree search. In H. Jaap van den Herik, Paolo Ciancarini, and H. H. L. M. (Jeroen) Donkers, editors, *Computers and Games*, pages 72–83, Berlin, Heidelberg, 2007. Springer Berlin Heidelberg.
- [2] Simeng Han, Hailey Schoelkopf, Yilun Zhao, Zhen-ting Qi, Martin Riddell, Wenfei Zhou, James Coady, David Peng, Yujie Qiao, Luke Benson, Lucy Sun, Alex Wardle-Solano, Hannah Szabo, Ekaterina Zubova, Matthew Burtell, Jonathan Fan, Yixin Liu, Brian Wong, Malcolm Sailor, Ansong Ni, Linyong Nan, Jungo Kasai, Tao Yu, Rui Zhang, Alexander R. Fabbri, Wojciech Kryscinski, Semih Yavuz, Ye Liu, Xi Victoria Lin, Shafiq Joty, Yingbo Zhou, Caiming Xiong, Rex Ying, Arman Cohan, and Dragomir Radev. Folio: Natural language reasoning with first-order logic, 2024.
- [3] Shen Nie, Fengqi Zhu, Zebin You, Xiaolu Zhang, Jingyang Ou, Jun Hu, Jun Zhou, Yankai Lin, Ji-Rong Wen, and Chongxuan Li. Large language diffusion models, 2025.
- [4] Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y. K. Li, Y. Wu, and Daya Guo. Deepseek-math: Pushing the limits of mathematical reasoning in open language models, 2024.